

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: path=r"C:\Users\jagad\Downloads\train_LZdllcl.csv"
df=pd.read_csv(path)
df
```

Out[3]:

	employee_id	department	region	education	gender	recruitment_channel	no_of_trainings
0	65438	Sales & Marketing	region_7	Master's & above	f	sourcing	1
1	65141	Operations	region_22	Bachelor's	m	other	1
2	7513	Sales & Marketing	region_19	Bachelor's	m	sourcing	1
3	2542	Sales & Marketing	region_23	Bachelor's	m	other	2
4	48945	Technology	region_26	Bachelor's	m	other	1
...
54803	3030	Technology	region_14	Bachelor's	m	sourcing	1
54804	74592	Operations	region_27	Master's & above	f	other	1
54805	13918	Analytics	region_1	Bachelor's	m	other	1
54806	13614	Sales & Marketing	region_9	NaN	m	sourcing	1
54807	51526	HR	region_22	Bachelor's	m	other	1

54808 rows × 14 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54808 entries, 0 to 54807
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   employee_id                          54808 non-null  int64
1   department                          54808 non-null  object
2   region                              54808 non-null  object
3   education                            52399 non-null  object
4   gender                              54808 non-null  object
5   recruitment_channel                  54808 non-null  object
6   no_of_trainings                      54808 non-null  int64
7   age                                 54808 non-null  int64
8   previous_year_rating                 50684 non-null  float64
9   length_of_service                   54808 non-null  int64
10  KPIs_met >80%                       54808 non-null  int64
11  awards_won?                         54808 non-null  int64
12  avg_training_score                   54808 non-null  int64
13  is_promoted                         54808 non-null  int64
dtypes: float64(1), int64(8), object(5)
memory usage: 5.9+ MB
```

```
In [5]: df.nunique()
```

```
Out[5]: employee_id      54808
department      9
region          34
education        3
gender           2
recruitment_channel  3
no_of_trainings  10
age             41
previous_year_rating  5
length_of_service  35
KPIs_met >80%      2
awards_won?        2
avg_training_score  61
is_promoted        2
dtype: int64
```

```
In [6]: # dropping the ID column
df.drop('employee_id',axis=1,inplace=True)
df
```

Out[6]:

	department	region	education	gender	recruitment_channel	no_of_trainings	age	previous
0	Sales & Marketing	region_7	Master's & above	f	sourcing	1	35	
1	Operations	region_22	Bachelor's	m	other	1	30	
2	Sales & Marketing	region_19	Bachelor's	m	sourcing	1	34	
3	Sales & Marketing	region_23	Bachelor's	m	other	2	39	
4	Technology	region_26	Bachelor's	m	other	1	45	
...
54803	Technology	region_14	Bachelor's	m	sourcing	1	48	
54804	Operations	region_27	Master's & above	f	other	1	37	
54805	Analytics	region_1	Bachelor's	m	other	1	27	
54806	Sales & Marketing	region_9	NaN	m	sourcing	1	29	
54807	HR	region_22	Bachelor's	m	other	1	27	

54808 rows × 13 columns

In [7]: `df.nunique()`

Out[7]:

department	9
region	34
education	3
gender	2
recruitment_channel	3
no_of_trainings	10
age	41
previous_year_rating	5
length_of_service	35
KPIs_met >80%	2
awards_won?	2
avg_training_score	61
is_promoted	2

dtype: int64

In [8]: `cat_cols=df.select_dtypes(include='object')`
`cat_cols.columns`

Out[8]: Index(['department', 'region', 'education', 'gender', 'recruitment_channel'], dtype='object')

In [9]: `num_cols=df.select_dtypes(exclude='object')`
`num_cols.columns`

Out[9]: Index(['no_of_trainings', 'age', 'previous_year_rating', 'length_of_service', 'KPIs_met >80%', 'awards_won?', 'avg_training_score', 'is_promoted'], dtype='object')

```
In [10]: df.isnull().sum()
```

```
Out[10]: department          0
region          0
education       2409
gender          0
recruitment_channel  0
no_of_trainings  0
age             0
previous_year_rating  4124
length_of_service  0
KPIs_met >80%      0
awards_won?      0
avg_training_score  0
is_promoted      0
dtype: int64
```

```
In [11]: df['education'].fillna(method='ffill',inplace=True)
median=df['previous_year_rating'].median()
print(median)
df['previous_year_rating'].fillna(median,inplace=True)
```

```
3.0
```

```
In [12]: df.isnull().sum()
```

```
Out[12]: department          0
region          0
education          0
gender          0
recruitment_channel  0
no_of_trainings  0
age             0
previous_year_rating  0
length_of_service  0
KPIs_met >80%      0
awards_won?      0
avg_training_score  0
is_promoted      0
dtype: int64
```

```
In [13]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
for i in cat_cols:
    df[i]=le.fit_transform(df[i])
df
```

Out[13]:

	department	region	education	gender	recruitment_channel	no_of_trainings	age	previous_y
0	7	31	2	0	2	1	35	
1	4	14	0	1	0	1	30	
2	7	10	0	1	2	1	34	
3	7	15	0	1	0	2	39	
4	8	18	0	1	0	1	45	
...
54803	8	5	0	1	2	1	48	
54804	4	19	2	0	0	1	37	
54805	0	0	0	1	0	1	27	
54806	7	33	0	1	2	1	29	
54807	2	14	0	1	0	1	27	

54808 rows × 13 columns

```
In [14]: # Data is ready
# splitting into X and y
X=df.drop('is_promoted',axis=1)
y=df['is_promoted']
```

```
In [15]: X.shape,y.shape
```

```
Out[15]: ((54808, 12), (54808,))
```

```
In [16]: # train_test_split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=42,test_size=0.2)
```

```
In [17]: X_train.shape,y_train.shape
```

```
Out[17]: ((43846, 12), (43846,))
```

```
In [18]: # Model Development with HyperParameter tuning Decision Tree
#metrics
from sklearn.metrics import accuracy_score,r2_score,\
    precision_score,recall_score,\
    f1_score,classification_report,\
    roc_auc_score,roc_curve,auc
from sklearn.metrics import confusion_matrix,ConfusionMatrixDisplay

from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV,cross_val_score
grid_tree=DecisionTreeClassifier() # base model
grid_tree
```

Out[18]: ▾ DecisionTreeClassifier
DecisionTreeClassifier()

In [19]: grid_tree.get_params()

Out[19]: {'ccp_alpha': 0.0,
'class_weight': None,
'criterion': 'gini',
'max_depth': None,
'max_features': None,
'max_leaf_nodes': None,
'min_impurity_decrease': 0.0,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'random_state': None,
'splitter': 'best'}

In [20]: *# creating a dictionary of hyper_parameters*
param_grid={'criterion':['gini','entropy'], #2
 'max_depth':[3,4,5,6,7,8], #6
 'min_samples_split':[2,3,4], #3
 'min_samples_leaf':[1,2,3,4], #4
 'random_state':[0,42,1234] } #3

param_grid

Out[20]: {'criterion': ['gini', 'entropy'],
'max_depth': [3, 4, 5, 6, 7, 8],
'min_samples_split': [2, 3, 4],
'min_samples_leaf': [1, 2, 3, 4],
'random_state': [0, 42, 1234]}

In [21]: *# now these parameters we apply in decision tree*
base model and params are ready now we apply GridSearchCV

grid_search= GridSearchCV(grid_tree, #base model
 param_grid, # parameters
 scoring='accuracy', # metrics
 cv=5,
 verbose=True)

grid_search

Out[21]: ▸ **GridSearchCV**
▸ estimator: DecisionTreeClassifier
 ▸ DecisionTreeClassifier

In [22]: import time
start=time.time()
grid_search.fit(X_train,y_train) # in dt without hyper parameters we fit directly in
end=time.time()
print('Total time taken is : ', end-start)

Fitting 5 folds for each of 432 candidates, totalling 2160 fits
Total time taken is : 194.02883791923523

In [23]: `grid_search.best_estimator_`

Out[23]: `DecisionTreeClassifier`
`DecisionTreeClassifier(criterion='entropy', max_depth=8, min_samples_leaf=2, random_state=0)`

In [24]: `grid_search.best_params_`

Out[24]: `{'criterion': 'entropy',
 'max_depth': 8,
 'min_samples_leaf': 2,
 'min_samples_split': 2,
 'random_state': 0}`

```
In [25]: from sklearn.tree import DecisionTreeClassifier
dtree=DecisionTreeClassifier(criterion='entropy',max_depth=8,min_samples_leaf=2,min_s
dtree.fit(X_train,y_train)

# model prediction
y_pred_dt=dtree.predict(X_test)
y_pred_dt

#metrics
from sklearn.metrics import accuracy_score,r2_score,\
    precision_score,recall_score,\
    f1_score,classification_report,\
    roc_auc_score,roc_curve, auc,\
    confusion_matrix,ConfusionMatrixDisplay

# confusion_matrix
cmt_dt = confusion_matrix(y_test,y_pred_dt)

# making confusion matrix as flat array
tn , fp ,fn ,tp = cmt_dt.ravel()

#accuracy_score,precision_score,Recall_score,f1 score
acc_dt=round(accuracy_score(y_test,y_pred_dt)*100,2)
precision_dt=round(precision_score(y_test,y_pred_dt)*100,2)
Recall_dt=round(recall_score(y_test,y_pred_dt)*100,2)
f1_dt=round(f1_score(y_test,y_pred_dt)*100,2)
print(f"Accuracy is = {acc_dt}%",)
print(f"F1 is = {f1_dt}%",)
print(f"Precision is = {precision_dt}%",)
print(f"Recall is = {Recall_dt}%",)

# instead of using seaborn we use ConfusionMatrixDisplay
#ConfusionMatrixDisplay(cmt_dt).plot()
#plt.grid(False)

dtree.predict_proba(X_test)
y_dt_pred_prob=dtree.predict_proba(X_test)[:,:1] # class-1 probabilities
fpr, tpr, threshold=roc_curve(y_test,y_dt_pred_prob)
plt.plot([0,1],[0,1],color="navy",lw=2,label="Random-Model")
plt.plot(fpr,tpr,color="darkorange",lw=2,label="Decision Tree Model")
plt.xlabel("<----- FPR ----->")
plt.ylabel("<----- TPR ----->")
plt.title("Receiver Operating Characteristics : ROC_AUC")
```

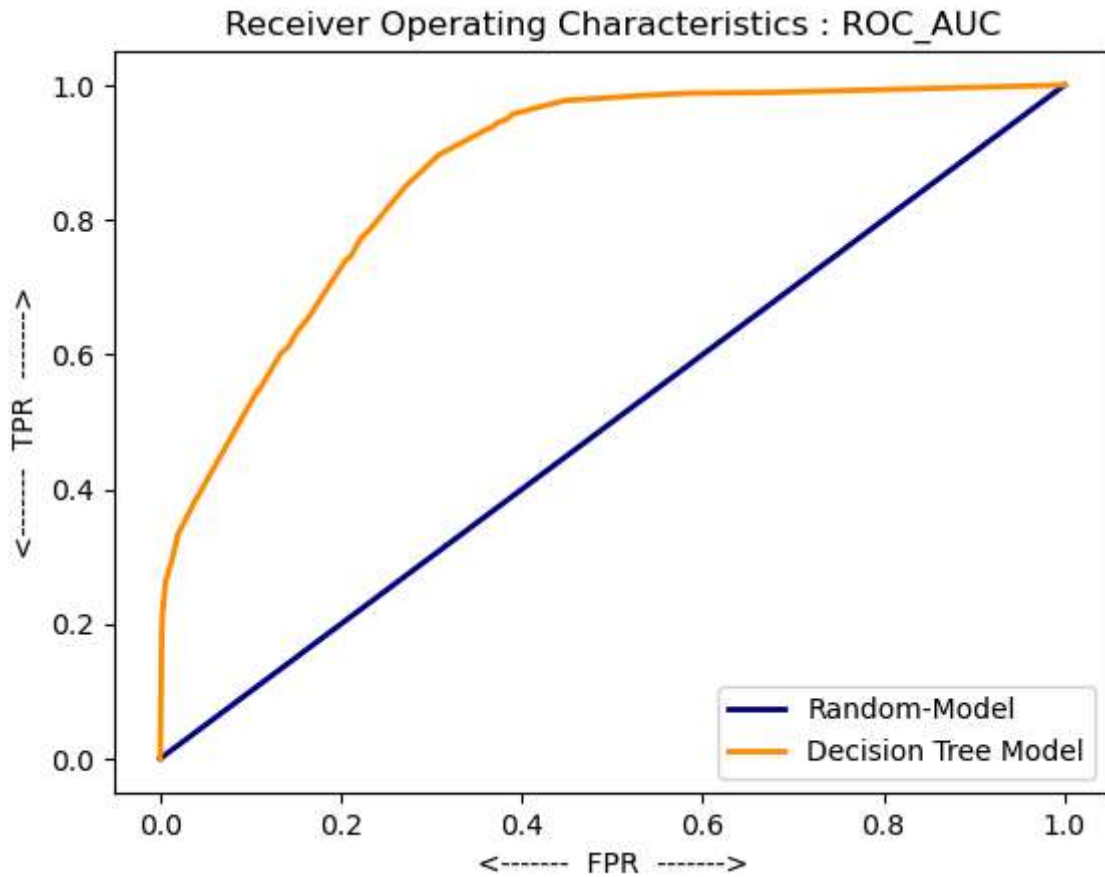
```
plt.legend()  
plt.show()  
AUC_dt=auc(fpr, tpr)  
print("Computed Area Under the Curve (AUC) ", AUC_dt)
```

Accuracy is = 93.36%

F1 is = 39.43%

Precision is = 80.61%

Recall is = 26.1%



Computed Area Under the Curve (AUC) 0.8702830705380374

```
In [26]: test=pd.read_csv(r"C:\Users\jagad\Downloads\test_2umaH9m.csv")  
test
```


Out[26]:

	employee_id	department	region	education	gender	recruitment_channel	no_of_trainings
0	8724	Technology	region_26	Bachelor's	m	sourcing	1
1	74430	HR	region_4	Bachelor's	f	other	1
2	72255	Sales & Marketing	region_13	Bachelor's	m	other	1
3	38562	Procurement	region_2	Bachelor's	f	other	3
4	64486	Finance	region_29	Bachelor's	m	sourcing	1
...
23485	53478	Legal	region_2	Below Secondary	m	sourcing	1
23486	25600	Technology	region_25	Bachelor's	m	sourcing	1
23487	45409	HR	region_16	Bachelor's	f	sourcing	1
23488	1186	Procurement	region_31	Bachelor's	m	sourcing	3
23489	5973	Technology	region_17	Master's & above	m	other	3

23490 rows × 13 columns

In [27]: `id=test['employee_id']`
`id`

Out[27]:

0	8724
1	74430
2	72255
3	38562
4	64486
...	...
23485	53478
23486	25600
23487	45409
23488	1186
23489	5973

Name: employee_id, Length: 23490, dtype: int64

In [28]: `test.drop('employee_id',axis=1,inplace=True)`

`test['education'].fillna(method='ffill',inplace=True)`
`median=test['previous_year_rating'].median()`
`print(median)`
`test['previous_year_rating'].fillna(median,inplace=True)`

3.0

In [29]: `test.isnull().sum()`

```
Out[29]: department      0
region      0
education    0
gender       0
recruitment_channel  0
no_of_trainings  0
age          0
previous_year_rating  0
length_of_service  0
KPIs_met >80%    0
awards_won?     0
avg_training_score  0
dtype: int64
```

```
In [30]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
for i in cat_cols:
    test[i]=le.fit_transform(test[i])
test
```

```
Out[30]:
```

	department	region	education	gender	recruitment_channel	no_of_trainings	age	previous_y
0	8	18	0	1	2	1	24	
1	2	28	0	0	0	1	31	
2	7	4	0	1	0	1	31	
3	5	11	0	0	0	3	31	
4	1	21	0	1	2	1	30	
...
23485	3	11	1	1	2	1	24	
23486	8	17	0	1	2	1	31	
23487	2	7	0	0	2	1	26	
23488	5	24	0	1	2	3	27	
23489	8	8	2	1	0	3	40	

23490 rows × 12 columns

```
In [31]: y_final_pred_dt=dtree.predict(test)
y_final_pred_dt
```

```
Out[31]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [32]: dff1=pd.DataFrame(zip(id,y_final_pred_dt),columns=['employee_id','is_promoted'])
dff1
```

Out[32]:

	employee_id	is_promoted
0	8724	0
1	74430	0
2	72255	0
3	38562	0
4	64486	0
...
23485	53478	0
23486	25600	0
23487	45409	0
23488	1186	0
23489	5973	0

23490 rows × 2 columns

```
In [33]: dff1.to_csv('final_submission.csv', index=False)
```

```
In [ ]:
```